

## 1. Image Classification

**Definition:** Assigns one or more labels to an entire image.

```
from torchvision import models, transforms
from PIL import Image
import torch

model = models.resnet18(pretrained=True).eval()
img = Image.open("image.jpg")
transform = transforms.Compose([
    transforms.Resize(256),
    transforms.CenterCrop(224),
    transforms.ToTensor()
])
input_tensor = transform(img).unsqueeze(0)
output = model(input_tensor)
print("Predicted class index:", output.argmax().item())
```

## 2. Object Detection

**Definition:** Identifies and localizes objects with bounding boxes.

```
import cv2

net = cv2.dnn.readNet("yolov3.weights", "yolov3.cfg")
img = cv2.imread("image.jpg")
blob = cv2.dnn.blobFromImage(img, 1/255, (416, 416), swapRB=True)
net.setInput(blob)
outputs = net.forward(net.getUnconnectedOutLayersNames())

for output in outputs:
    for detection in output:
        scores = detection[5:]
        class_id = scores.argmax()
        confidence = scores[class_id]
        if confidence > 0.5:
            print("Detected class ID:", class_id)
```

### 3. Semantic Segmentation

**Definition:** Classifies each pixel into a semantic category.

```
import torchvision
from PIL import Image
import torch

model = torchvision.models.segmentation.deeplabv3_resnet50(pretrained=True).eval()
img = Image.open("image.jpg").convert("RGB")
transform = torchvision.transforms.Compose([
    torchvision.transforms.Resize((520, 520)),
    torchvision.transforms.ToTensor()
])
input_tensor = transform(img).unsqueeze(0)
output = model(input_tensor)[“out”].argmax(1).squeeze().numpy()
print("Segmentation map shape:", output.shape)
```

### 4. Instance Segmentation

**Definition:** Segments and distinguishes individual object instances.

```
import torchvision
from PIL import Image
import torch
from torchvision.transforms import functional as F

model = torchvision.models.detection.maskrcnn_resnet50_fpn(pretrained=True).eval()
img = Image.open("image.jpg").convert("RGB")
input_tensor = F.to_tensor(img).unsqueeze(0)
output = model(input_tensor)[0]

for i, mask in enumerate(output[‘masks’]):
    print(f"Instance {i} mask shape:", mask.shape)
```

## 5. Panoptic Segmentation

**Definition:** Combines semantic and instance segmentation.

```
# Panoptic segmentation is available via Detectron2.  
# Here's a basic example using Hugging Face Transformers:  
  
from transformers import AutoImageProcessor, AutoModelForPanopticSegmentation  
from PIL import Image  
import torch  
  
img = Image.open("image.jpg")  
processor = AutoImageProcessor.from_pretrained("facebook/detectron2")  
model = AutoModelForPanopticSegmentation.from_pretrained("facebook/detectron2")  
  
inputs = processor(images=img, return_tensors="pt")  
outputs = model(**inputs)  
print("Panoptic segmentation output:", outputs.keys())
```

## 6. Keypoint Detection

**Definition:** Detects anatomical or structural points (e.g., joints).

```
import mediapipe as mp  
import cv2  
  
pose = mp.solutions.pose.Pose()  
img = cv2.imread("person.jpg")  
img_rgb = cv2.cvtColor(img, cv2.COLOR_BGR2RGB)  
results = pose.process(img_rgb)  
  
if results.pose_landmarks:  
    for lm in results.pose_landmarks.landmark:  
        print("Keypoint:", lm.x, lm.y)
```

## 7. Pose Estimation

**Definition:** Predicts human or object poses via keypoints.

```
import cv2
import mediapipe as mp

# Initialize MediaPipe pose class and drawing utilities
mp_pose = mp.solutions.pose
mp_drawing = mp.solutions.drawing_utils

# Load video or webcam
cap = cv2.VideoCapture(0) # Use 0 for webcam

with mp_pose.Pose(static_image_mode=False,
                   model_complexity=1,
                   enable_segmentation=False,
                   min_detection_confidence=0.5,
                   min_tracking_confidence=0.5) as pose:

    while cap.isOpened():
        success, frame = cap.read()
        if not success:
            break

        # Convert the image to RGB
        image = cv2.cvtColor(frame, cv2.COLOR_BGR2RGB)
        image.flags.writeable = False

        # Process the image and detect pose
        results = pose.process(image)

        # Draw pose landmarks on the image
        image.flags.writeable = True
        image = cv2.cvtColor(image, cv2.COLOR_RGB2BGR)

        if results.pose_landmarks:
            mp_drawing.draw_landmarks(
                image, results.pose_landmarks, mp_pose.POSE_CONNECTIONS)

        # Display the result
        cv2.imshow('Pose Estimation', image)

        if cv2.waitKey(5) & 0xFF == 27:
            break

cap.release()
cv2.destroyAllWindows()
```

## 8. Facial Recognition

**Definition:** Identifies or verifies individual faces.

```
import face_recognition

img = face_recognition.load_image_file("face.jpg")
encodings = face_recognition.face_encodings(img)
print("Face encoding vector:", encodings[0])
```

## 9. Image Generation (Synthesis)

**Definition:** Generates realistic or stylized images from input.

```
from diffusers import StableDiffusionPipeline

pipe = StableDiffusionPipeline.from_pretrained("runwayml/stable-diffusion-v1-5").to("cuda")
image = pipe("a futuristic city at sunset").images[0]
image.save("generated.png")
```

## 10. Image-to-Image Translation

**Definition:** Converts one image domain to another.

```
# Example using Pix2Pix via Hugging Face Diffusers

from diffusers import Pix2PixZeroPipeline
pipe = Pix2PixZeroPipeline.from_pretrained("timbrooks/pix2pix-zero-portrait").to("cuda")
image = pipe(prompt="make it look like a cartoon", image="portrait.jpg").images[0]
image.save("translated.png")
```

## 11. Super Resolution

**Definition:** Enhances image resolution.

```
from realesrgan import RealESRGAN
from PIL import Image

model = RealESRGAN("cuda", scale=2)
model.load_weights("RealESRGAN_x2.pth")
img = Image.open("low_res.jpg")
sr_img = model.predict(img)
sr_img.save("high_res.jpg")
```

## 12. Image Denoising

**Definition:** Removes noise from images.

```
import cv2
img = cv2.imread("noisy.jpg")
denoised = cv2.fastNlMeansDenoisingColored(img, None, 10, 10, 7, 21)
cv2.imwrite("denoised.jpg", denoised)
```

## 13. Image Inpainting

**Definition:** Fills in missing regions.

```
import cv2
img = cv2.imread("damaged.jpg")
mask = cv2.imread("mask.jpg", 0) # white where missing
inpainted = cv2.inpaint(img, mask, 3, cv2.INPAINT_TELEA)
cv2.imwrite("inpainted.jpg", inpainted)
```

## 14. Image Captioning

**Definition:** Generates natural language descriptions of images.

```
from transformers import BlipProcessor, BlipForConditionalGeneration
from PIL import Image

processor = BlipProcessor.from_pretrained("Salesforce/blip-image-captioning-base")
model = BlipForConditionalGeneration.from_pretrained("Salesforce/blip-image-captioning-
base")

img = Image.open("image.jpg")
inputs = processor(images=img, return_tensors="pt")
caption = model.generate(**inputs)
print("Caption:", processor.decode(caption[0], skip_special_tokens=True))
```

## 15. Visual Question Answering (VQA)

**Definition:** Answers questions based on image content.

```
from transformers import ViltProcessor, ViltForQuestionAnswering
from PIL import Image

processor = ViltProcessor.from_pretrained
```

## 16. Depth Estimation

**Definition:** Infers depth from single or stereo images.

```
from transformers import DPTFeatureExtractor, DPTForDepthEstimation
from PIL import Image
import torch

img = Image.open("scene.jpg")
extractor = DPTFeatureExtractor.from_pretrained("Intel/dpt-large")
model = DPTForDepthEstimation.from_pretrained("Intel/dpt-large")

inputs = extractor(images=img, return_tensors="pt")
with torch.no_grad():
    depth = model(**inputs).predicted_depth
print("Depth map shape:", depth.shape)
```

## 17. 3D Reconstruction

**Definition:** Rebuilds 3D models from 2D images or video.

💡 **Note:** This task typically uses external tools like COLMAP or [OpenMVG]. Here's how to run COLMAP from Python using `subprocess`

```
import subprocess

# Paths to your image folder and output workspace
image_path = "./images"
workspace_path = "./output"

# Run COLMAP automatic reconstruction
subprocess.run([
    "colmap", "automatic_reconstructor",
    "--image_path", image_path,
    "--workspace_path", workspace_path,
    "--data_type", "image",
    "--use_gpu", "1"
])
# This will generate a sparse and dense 3D model from an image set.
```

## 18. Optical Flow Estimation

**Definition:** Tracks motion between video frames.

```
import cv2

prev = cv2.imread("frame1.jpg", cv2.IMREAD_GRAYSCALE)
next = cv2.imread("frame2.jpg", cv2.IMREAD_GRAYSCALE)
flow = cv2.calcOpticalFlowFarneback(prev, next, None, 0.5, 3, 15, 3, 5, 1.2, 0)
print("Optical flow shape:", flow.shape)
```

## 19. Video Classification

**Definition:** Assigns labels to video clips.

```
# Example using Hugging Face TimeSformer:  
from transformers import TimesformerProcessor, TimesformerForVideoClassification  
import torch  
from decord import VideoReader, cpu  
  
video = VideoReader("video.mp4", ctx=cpu())  
frames = [video[i].asnumpy() for i in range(0, len(video), 4)]  
  
processor = TimesformerProcessor.from_pretrained("facebook/timesformer-base-finetuned-k400")  
model = TimesformerForVideoClassification.from_pretrained("facebook/timesformer-base-finetuned-k400")  
  
inputs = processor(videos=frames, return_tensors="pt")  
outputs = model(**inputs)  
label = model.config.label2id[outputs.logits.argmax().item()]  
print("Predicted label:", label)
```

## 20. Action Recognition

**Definition:** Identifies human actions from video.

### Using PyTorchVideo's pretrained SlowFast model:

```
from pytorchvideo.models.hub import slowfast_r50
import torch
import torchvision.transforms as transforms
import numpy as np
import cv2

# Load pretrained model
model = slowfast_r50(pretrained=True).eval()

# Load video and extract frames
cap = cv2.VideoCapture("action.mp4")
frames = []
while len(frames) < 32:
    ret, frame = cap.read()
    if not ret: break
    frame = cv2.resize(frame, (256, 256))
    frames.append(frame)
cap.release()

# Preprocess frames
transform = transforms.Compose([
    transforms.ToTensor(),
    transforms.Normalize([0.45, 0.45, 0.45], [0.225, 0.225, 0.225])
])
video_tensor = torch.stack([transform(f) for f in frames]).permute(1, 0, 2,
3).unsqueeze(0)

# Run model
with torch.no_grad():
    output = model(video_tensor)
    predicted_class = output.argmax().item()
print("Predicted action class index:", predicted_class)
```

## 21. Object Tracking

**Definition:** Follows object location over time in video.

```
import cv2

tracker = cv2.TrackerCSRT_create()
video = cv2.VideoCapture("video.mp4")
ret, frame = video.read()
bbox = cv2.selectROI("Tracking", frame, False)
tracker.init(frame, bbox)

while True:
    ret, frame = video.read()
    if not ret: break
    success, box = tracker.update(frame)
    if success:
        x, y, w, h = [int(v) for v in box]
        cv2.rectangle(frame, (x, y), (x+w, y+h), (0,255,0), 2)
    cv2.imshow("Tracking", frame)
    if cv2.waitKey(1) == 27: break
video.release()
cv2.destroyAllWindows()
```

## 22. Multimodal Learning

**Definition:** Combines vision with language or audio.

```
from transformers import CLIPProcessor, CLIPModel
from PIL import Image
import torch

model = CLIPModel.from_pretrained("openai/clip-vit-base-patch32")
processor = CLIPProcessor.from_pretrained("openai/clip-vit-base-patch32")

img = Image.open("image.jpg")
text = ["a dog", "a cat", "a car"]

inputs = processor(text=text, images=img, return_tensors="pt", padding=True)
outputs = model(**inputs)
probs = outputs.logits_per_image.softmax(dim=1)
print("Text match probabilities:", probs)
```

## 23. Scene Text Recognition (OCR)

**Definition:** Detects and reads text from images.

```
import pytesseract
import cv2

img = cv2.imread("text_image.jpg")
text = pytesseract.image_to_string(img)
print("Extracted text:", text)
```

## 24. Image Retrieval

**Definition:** Finds visually similar images.

```
from PIL import Image
from transformers import CLIPProcessor, CLIPModel
import torch

model = CLIPModel.from_pretrained("openai/clip-vit-base-patch32")
processor = CLIPProcessor.from_pretrained("openai/clip-vit-base-patch32")

def load_image(path): return Image.open(path).convert("RGB")
image_paths = ["query.jpg", "img1.jpg", "img2.jpg"]
images = [load_image(p) for p in image_paths]

inputs = processor(images=images, return_tensors="pt", padding=True)
with torch.no_grad():
    embeddings = model.get_image_features(**inputs)
embeddings = embeddings / embeddings.norm(dim=-1, keepdim=True)

query = embeddings[0]
similarities = torch.matmul(embeddings[1:], query.unsqueeze(0).T).squeeze()
for i in similarities.argsort(descending=True):
    print(f"Match: {image_paths[i+1]}, Score: {similarities[i].item():.4f}")
```

## 25. Domain Adaptation

**Definition:** Applies vision models to new but similar domains.

### Using a simplified Domain-Adversarial Neural Network (DANN) setup:

```
import torch
import torch.nn as nn
from torch.autograd import Function

# Gradient reversal layer
class GradReverse(Function):
    @staticmethod
    def forward(ctx, x, lambd): ctx.lambd = lambd; return x.view_as(x)
    @staticmethod
    def backward(ctx, grad_output): return -ctx.lambd * grad_output, None

# Feature extractor
class FeatureExtractor(nn.Module):
    def __init__(self): super().__init__()
    self.conv = nn.Sequential(nn.Conv2d(3, 64, 3), nn.ReLU(), nn.Flatten())

    def forward(self, x): return self.conv(x)

# Label predictor
class LabelPredictor(nn.Module):
    def __init__(self): super().__init__()
    self.fc = nn.Linear(64 * 30 * 30, 10)

    def forward(self, x): return self.fc(x)

# Domain classifier
class DomainClassifier(nn.Module):
    def __init__(self): super().__init__()
    self.fc = nn.Linear(64 * 30 * 30, 2)

    def forward(self, x, lambd): x = GradReverse.apply(x, lambd); return self.fc(x)
    # You'd train this with source and target domain batches, optimizing both label and
    domain losses.
```

## 26. Zero-shot Learning (ZSL)

**Definition:** Recognizes novel classes not seen during training.

```
from transformers import CLIPProcessor, CLIPModel
from PIL import Image

model = CLIPModel.from_pretrained("openai/clip-vit-base-patch32")
processor = CLIPProcessor.from_pretrained("openai/clip-vit-base-patch32")

img = Image.open("image.jpg")
labels = ["a zebra", "a giraffe", "a lion"]
inputs = processor(text=labels, images=img, return_tensors="pt", padding=True)
outputs = model(**inputs)
pred = labels[outputs.logits_per_image.argmax()]
print("Zero-shot prediction:", pred)
```

## 27. Few-shot Learning

**Definition:** Learns from very few labeled examples.

📦 **Using Prototypical Networks (simplified):**

```
import torch
import torch.nn.functional as F

# Support set: 5 examples per class
support = torch.randn(10, 64) # 2 classes x 5 samples
support_labels = torch.tensor([0]*5 + [1]*5)

# Query set: unlabeled examples
query = torch.randn(3, 64)

# Compute prototypes
prototypes = torch.stack([
    support[support_labels == c].mean(0) for c in torch.unique(support_labels)
])

# Compute distances and classify
dists = torch.cdist(query, prototypes)
preds = dists.argmin(dim=1)
print("Predicted classes:", preds.tolist())
```

## 28. Image Style Transfer

**Definition:** Applies artistic style from one image to another.

```
import torch
import torchvision.transforms as transforms
from PIL import Image
from torchvision.models import vgg19

# Load content and style images
content = Image.open("content.jpg")
style = Image.open("style.jpg")
transform = transforms.Compose([transforms.Resize((256, 256)), transforms.ToTensor()])
content_tensor = transform(content).unsqueeze(0)
style_tensor = transform(style).unsqueeze(0)

# Use VGG features and optimize content image to match style
# Requires full neural style transfer pipeline (omitted for brevity)
```

## 29. Image Matting

**Definition:** Extracts foreground object with fine boundaries.

📦 **Using MODNet via ONNX Runtime:**

```
import onnxruntime
import cv2
import numpy as np

# Load MODNet ONNX model
session = onnxruntime.InferenceSession("modnet.onnx")

# Preprocess image
img = cv2.imread("portrait.jpg")
img = cv2.resize(img, (512, 512))
input_blob = img.transpose(2, 0, 1).astype(np.float32) / 255.0
input_blob = np.expand_dims(input_blob, axis=0)

# Run inference
outputs = session.run(None, {"input": input_blob})
alpha = outputs[0][0][0] # Alpha matte

# Apply matte
foreground = (img * alpha[..., None]).astype(np.uint8)
cv2.imwrite("matte.png", foreground)
```

## 30. Image Quality Assessment

**Definition:** Evaluates perceptual quality of images.

📦 Using BRISQUE via `piq` (PyTorch Image Quality):

```
import torch
import piq
from PIL import Image
from torchvision.transforms import ToTensor

img = Image.open("test.jpg").convert("RGB")
img_tensor = ToTensor()(img).unsqueeze(0)

score = piq.brisque(img_tensor)
print("BRISQUE score:", score.item())
```