

1 Generative Models for Vision Synthesis

Definition

Generative models learn patterns in data to create new, synthetic content. In vision: they generate images, transform styles, or reconstruct missing data.

◆ GANs (Generative Adversarial Networks)

- **Concept:** Competing neural networks (generator vs discriminator) generate highly realistic samples.
- **Code Snippet: Basic DCGAN Generator**

```
class Generator(nn.Module):  
  
    def __init__(self):  
        super().__init__()  
  
        self.model = nn.Sequential(  
  
            nn.ConvTranspose2d(100, 64, 4, 1, 0),  
            nn.ReLU(),  
            nn.ConvTranspose2d(64, 1, 4, 2, 1),  
            nn.Tanh()  
        )  
  
    def forward(self, x):  
        return self.model(x)
```

◆ VAEs (Variational Autoencoders)

- **Concept:** Learn latent space representations and reconstruct images while allowing smooth interpolation.
- **Use Case:** Facial morphing, anomaly detection

Project: Create a synthetic fashion/image dataset using GAN or VAE

Tools: PyTorch, torchvision, numpy, TensorBoard

2 Vision Transformers (ViTs) and Attention Mechanisms

Definition

Transformers model relationships between image patches via self-attention, allowing long-range dependencies in image understanding.

- ◆ **Vision Transformer (ViT)**

- **Code Snippet: Load with timm**

```
import timm  
model = timm.create_model("vit_base_patch16_224", pretrained=True)
```

- **Project:** Compare ViT vs. CNN accuracy on ImageNet subsets

-  Tools: PyTorch, timm, HuggingFace datasets

3 Multimodal Learning: Vision + Text/Audio

Definition

Models that combine multiple input types—such as images and text—to answer queries, generate captions, or perform scene reasoning.

- ◆ **VQA (Visual Question Answering)**

- **Concept:** System answers questions based on image input.
 - **Architecture:** CNN (image features) + Transformer or RNN (text query)
 - **Code Snippet (text tokenization + image embedding)**

```
from transformers import BertTokenizer  
  
tokenizer = BertTokenizer.from_pretrained("bert-base-uncased")  
  
tokens = tokenizer("What color is the cat?", return_tensors="pt")
```

- **Project:** Build a VQA model using Flickr or GQA datasets

-  Tools: PyTorch, HuggingFace, torchvision, OpenCLIP

4 Ethical AI, Fairness, Bias Mitigation

Definition

Ensure models avoid discriminatory behavior by addressing dataset imbalance, harmful outputs, and inaccessible designs.

- ◆ Concepts

- **Bias sources:** Label imbalance, cultural representation, resolution disparities
- **Techniques:** Balanced sampling, adversarial fairness training, explainability

Project: Audit CV model for demographic bias on face detection

 Tools: IBM AI Fairness 360, Captum, SHAP, custom dataset analysis

5 Publishing Research in Top-tier Journals

Goals

Train students to structure research papers for conferences (CVPR, ICCV, NeurIPS). Emphasis on reproducibility, open science, and visual storytelling.

- ◆ Module Includes:

- Writing abstracts, introductions, and impact statements
- Crafting method illustrations and diagrams
- Using LaTeX for academic formatting
- Open-source release documentation

Project: Write and document a reproducible image synthesis pipeline

 Tools: Overleaf, GitHub Pages, Jupyter Book, wandb

Markdown Template Example of Vision Transformer Module

Overview

Vision Transformers segment images into patches and process them using self-attention layers. This enables contextual understanding across large spatial regions.

Dependencies

```
pip install timm torch torchvision
```

Basic Usage

```
import timm  
  
model = timm.create_model('vit_base_patch16_224', pretrained=True)
```

Visual Diagram

[Image] → [Patch Embedding] → [Transformer Encoder] → [Classification Head]

Notes

ViTs require large datasets and compute but offer state-of-the-art results on classification tasks.

Code Examples Galore

| Task | Model | Library |

|-----|-----|-----|

| Image synthesis | DCGAN, VAE | PyTorch |

| Classification | ViT, SwinTransformer | timm |

| VQA | ViLT, CLIP + BERT | HuggingFace |

| Bias analysis | FaceNet + fairness metrics | Captum, AI Fairness 360 |

| Paper figures | Custom data plotter | matplotlib, LaTeX PGFPlots |

■ Generative Models (GANs & VAEs) for Image Synthesis

```
import torch  
  
import torch.nn as nn  
  
import matplotlib.pyplot as plt  
  
  
  
# Generator (DCGAN-style)  
  
class Generator(nn.Module):  
  
    def __init__(self):  
  
        super().__init__()  
  
        self.model = nn.Sequential(  
  
            nn.ConvTranspose2d(100, 64, 4, 1, 0),  
  
            nn.BatchNorm2d(64),  
  
            nn.ReLU(True),  
  
            nn.ConvTranspose2d(64, 1, 4, 2, 1),  
  
            nn.Tanh()  
  
        )  
  
        def forward(self, x):  
  
            return self.model(x)  
  
  
  
# Sample noise & generate image  
  
noise = torch.randn(1, 100, 1, 1)  
  
gen = Generator()  
  
img = gen(noise).detach().squeeze()  
  
plt.imshow(img.numpy(), cmap="gray")  
  
plt.title("Synthetic Image from GAN Generator")  
  
plt.axis("off")  
  
plt.show()
```

📘 Vision Transformers & Attention with timm:

```
!pip install timm  
  
import timm  
  
import torch  
  
  
# Load pretrained ViT  
  
model = timm.create_model('vit_base_patch16_224', pretrained=True)  
  
# Summary of architecture  
  
print(model)
```

Use hooks or forward pass to extract embeddings later

🔍 *Optional: Add a visualization cell for patch embeddings over image using OpenCV or matplotlib.*

📘 Multimodal VQA System

```
# Visual Question Answering Prototype  
  
from transformers import BertTokenizer, BertModel  
  
import torch  
  
import torchvision.models as models  
  
  
# Text tokenization  
  
tokenizer = BertTokenizer.from_pretrained("bert-base-uncased")  
  
tokens = tokenizer("What color is the cat?", return_tensors="pt")  
  
# Image feature extractor  
  
img_model = models.resnet18(pretrained=True)  
  
img_model.fc = torch.nn.Identity() # Remove classifier  
  
# Example image tensor (load with torchvision or PIL)  
  
# img_tensor = preprocess("cat.jpg")
```

```
# Combine for fusion later: [Text tokens] + [Image features]
```

 Extend this by fusing modalities for end-to-end prediction.

Ethics and Bias Auditing > Bias Audit: Face Detection Example

```
from aif360.datasets import BinaryLabelDataset  
  
from aif360.metrics import BinaryLabelDatasetMetric  
  
  
# Simulated dataset load  
  
data = BinaryLabelDataset(  
  
    favorable_label=1, unfavorable_label=0,  
  
    df=your_dataframe, label_names=['prediction'], protected_attribute_names=['ethnicity'])  
  
)  
  
  
metric = BinaryLabelDatasetMetric(data, privileged_groups=[{'ethnicity': 1}],  
    unprivileged_groups=[{'ethnicity': 0}])  
  
print("Disparate impact:", metric.disparate_impact())
```

💡 You can analyze how fairness shifts with dataset balancing or adversarial correction.

📘 Research & Reproducibility

```
# Publishing Pipeline: Figure Generation & Logging
```

```
import matplotlib.pyplot as plt
```

```
import wandb
```

```
wandb.init(project="image-synthesis-pipeline")
```

```
loss_vals = [0.5, 0.4, 0.3, 0.2]
```

```
epochs = list(range(1, 5))
```

```
plt.plot(epochs, loss_vals)
```

```
plt.title("Training Loss Over Epochs")
```

```
plt.xlabel("Epoch")
```

```
plt.ylabel("Loss")
```

```
plt.grid(True)
```

```
plt.show()
```

```
# Log to wandb
```

```
wandb.log({"loss_chart": plt})
```

💡 GAN Architecture (DCGAN-style)

With `matplotlib`

```
import matplotlib.pyplot as plt

import matplotlib.patches as patches

fig, ax = plt.subplots(figsize=(10, 2))

# Generator

ax.add_patch(patches.Rectangle((0.1, 0.5), 0.3, 0.4, color='skyblue', label="Generator"))

ax.text(0.25, 0.65, "Generator", ha='center')

# Discriminator

ax.add_patch(patches.Rectangle((0.6, 0.5), 0.3, 0.4, color='salmon'))

ax.text(0.75, 0.65, "Discriminator", ha='center')

# Arrows

ax.annotate("", xy=(0.4, 0.7), xytext=(0.6, 0.7), arrowprops=dict(arrowstyle="->"))

ax.text(0.5, 0.75, "Synthetic Image", ha='center')

ax.annotate("", xy=(0.9, 0.7), xytext=(0.6, 0.7), arrowprops=dict(arrowstyle="->"))

ax.text(0.75, 0.85, "Adversarial Feedback", ha='center')

ax.axis('off')

plt.title("GAN Architecture Flow")

plt.show()
```

2. Vision Transformer Flow

With `plotly.graph_objects`

```
import plotly.graph_objects as go

fig = go.Figure()

steps = ["Image", "Patchify", "Embedding", "Self-Attention", "Classification"]

for i, step in enumerate(steps):
    fig.add_shape(type="rect", x0=i, x1=i+0.8, y0=0, y1=0.5,
                  line=dict(color="black"), fillcolor="lightgreen")
    fig.add_annotation(x=i+0.4, y=0.25, text=step, showarrow=False)

# Arrows
for i in range(len(steps)-1):
    fig.add_annotation(x=i+0.8, y=0.25, ax=i+0.9, ay=0.25,
                        arrowhead=2, arrowsize=1.5)

fig.update_layout(height=200, width=800, showlegend=False,
                  title="Vision Transformer Flow", plot_bgcolor="white",
                  margin=dict(l=20, r=20, t=40, b=20))

fig.show()
```

3. Multimodal Fusion Diagram

With graphviz

```
from graphviz import Digraph
```

```
dot = Digraph()
```

```
dot.node('I', 'Image → CNN')
```

```
dot.node('Q', 'Question → BERT')
```

```
dot.node('F', 'Fusion Layer')
```

```
dot.node('A', 'Answer Prediction')
```

```
dot.edge('I', 'F')
```

```
dot.edge('Q', 'F')
```

```
dot.edge('F', 'A')
```

```
dot.render('multimodal_fusion', format='png', cleanup=True)
```

4. Fairness Audit Flow

With `graphviz`

```
flow = Digraph()  
  
flow.node("D", "Image Dataset")  
  
flow.node("B", "Bias Detection")  
  
flow.node("M", "Metrics Report")  
  
flow.node("S", "Mitigation Strategy")  
  
flow.edge("D", "B")  
  
flow.edge("B", "M")  
  
flow.edge("M", "S")  
  
flow.render("fairness_flow", format="png", cleanup=True)
```

5. Paper Composition Map

With `matplotlib` Text Annotations

```
fig, ax = plt.subplots(figsize=(10, 2))  
  
sections = ["Title", "Abstract", "Introduction", "Methodology", "Experiments", "Conclusion"]  
  
for i, section in enumerate(sections):  
  
    ax.text(i, 0.5, section, fontsize=12, ha='center', bbox=dict(boxstyle="round",  
    facecolor="lightyellow"))  
  
    if i < len(sections) - 1:  
  
        ax.annotate("", xy=(i+0.5, 0.5), xytext=(i+0.9, 0.5), arrowprops=dict(arrowstyle="->"))  
  
ax.axis('off')  
  
plt.title("Research Paper Composition Flow")  
  
plt.show()
```

झूँझुका विकास

प्रश्न उत्तर सेलेक्टर विकास

```
from ipywidgets import interact, Dropdown  
import matplotlib.pyplot as plt  
  
questions = Dropdown(  
    options=["What is on the table?", "How many dogs?", "Is it daytime?"],  
    description='Question:',  
)  
  
def predict_answer(q):  
  
    img = plt.imread("sample.jpg")  
    plt.imshow(img)  
    plt.title(f"Predicted Answer: {q} → 'Bottle'")  
    plt.axis('off')  
    plt.show()  
  
interact(predict_answer, q=questions)
```

विचारना स्लायडर

```
from ipywidgets import FloatSlider  
import matplotlib.pyplot as plt  
import numpy as np  
  
def audit(threshold):  
  
    disparity = max(0, 1 - threshold)  
  
    plt.bar(['Privileged', 'Unprivileged'], [1.0, 1.0 - disparity], color=['green', 'red'])  
    plt.title(f"Fairness Audit @ Bias Threshold: {threshold}")  
    plt.ylabel("Acceptance Rate")  
    plt.show()  
  
interact(audit, threshold=FloatSlider(min=0, max=1, step=0.1, value=0.5))
```

⌚ Segmentation Overlay Widget

```
from ipywidgets import FloatSlider
import numpy as np
import matplotlib.pyplot as plt

def overlay(alpha=0.4):
    image = plt.imread("dashcam.jpg")
    mask = np.load("lane_mask.npy")
    plt.imshow(image)
    plt.imshow(mask, cmap='Reds', alpha=alpha)
    plt.title("Lane Segmentation Overlay")
    plt.axis('off')
    plt.show()

interact(overlay, alpha=FloatSlider(min=0.0, max=1.0, step=0.05, value=0.4))
```

Interactive Widgets for Simulation

Using `ipywidgets`, we can simulate model behavior and user input controls directly in notebooks.

GAN: Adjust Noise Vector Dimensionality

```
from ipywidgets import interact  
  
import torch  
  
import matplotlib.pyplot as plt  
  
  
def generate_image(noise_dim=100):  
  
    noise = torch.randn(1, noise_dim, 1, 1)  
  
    # Placeholder generator logic  
  
    img = noise.squeeze().detach().numpy()[:28, :28] # Fake output  
  
    plt.imshow(img, cmap='gray')  
  
    plt.title(f"GAN Output (dim={noise_dim})")  
  
    plt.axis('off')  
  
    plt.show()  
  
  
interact(generate_image, noise_dim=(10, 200, 10))
```

🧠 Transformer Attention Head Selector

```
from ipywidgets import Dropdown
import matplotlib.pyplot as plt
import numpy as np

attention_heads = Dropdown(
    options=['Head 1', 'Head 2', 'Head 3'],
    value='Head 1',
    description='Attention Head:',
)

def plot_attention(head):
    attn_map = np.random.rand(8, 8) # Fake 8x8 patch attention
    plt.imshow(attn_map, cmap='viridis')
    plt.title(f"Attention Map: {head}")
    plt.colorbar()
    plt.show()

attention_heads.observe(lambda change: plot_attention(change['new']), names='value')
plot_attention(attention_heads.value)
```

⚙️ 1. Real-Time Widget ↔ Pretrained Model Link

🎮 Example: VQA Widget with Inference Logic

```
import torch

from transformers import BlipProcessor, BlipForQuestionAnswering

from PIL import Image

from ipywidgets import interact, Dropdown

# Load pretrained BLIP VQA model

processor = BlipProcessor.from_pretrained("Salesforce/blip-vqa-base")

model = BlipForQuestionAnswering.from_pretrained("Salesforce/blip-vqa-base")

def vqa_inference(image_path, question):

    image = Image.open(image_path).convert("RGB")

    inputs = processor(image, question, return_tensors="pt")

    out = model.generate(**inputs)

    answer = processor.decode(out[0], skip_special_tokens=True)

    return answer

# Interactive widget

def ask(question):

    answer = vqa_inference("vqa_sample.jpg", question)

    print(f"Answer: {answer}")

    interact(ask, question=Dropdown(options=[

        "What is the person doing?",

        "Is the dog sitting?",

        "What color is the backpack?"

    ], description="Ask:"))
```

💡 This links directly to a pretrained model with instant answers. You can extend this to other models (YOLOv8, CLIP, segmentation networks) using similar logic.

2. FastAPI Backend for Deployment

Create a RESTful API that serves predictions to frontend widgets or external apps.

Install Dependencies

```
pip install fastapi uvicorn pillow transformers torch torchvision
```

File Structure

```
cv-api/
```

```
|
```

```
|   └── app.py
```

```
|   └── models/
```

```
|       └── vqa_model.py
```

```
|       └── segmentation_model.py
```

```
|   └── static/
```

```
|       └── sample.jpg
```

VQA Endpoint (app.py)

```
from fastapi import FastAPI, UploadFile, Form  
  
from models.vqa_model import answer_question  
  
app = FastAPI()  
  
@app.post("/vqa")  
  
async def vqa(image: UploadFile, question: str = Form(...)):  
  
    img_bytes = await image.read()  
  
    result = answer_question(img_bytes, question)  
  
    return {"answer": result}
```

Model Logic (models/vqa_model.py)

```
from transformers import BlipProcessor, BlipForQuestionAnswering  
  
from PIL import Image  
  
import io
```

```
processor = BlipProcessor.from_pretrained("Salesforce/blip-vqa-base")
model = BlipForQuestionAnswering.from_pretrained("Salesforce/blip-vqa-base")
```

```
def answer_question(image_bytes, question):
    image = Image.open(io.BytesIO(image_bytes)).convert("RGB")
    inputs = processor(image, question, return_tensors="pt")
    out = model.generate(**inputs)
    return processor.decode(out[0], skip_special_tokens=True)
```

🌐 Deployment & Demo Connection

- Run locally:
uvicorn app:app --reload
- To connect a widget frontend:
- Use JavaScript (`fetch()` or `axios`)
- Or send requests via Python notebooks:

```
import requests
```

```
files = {'image': open('vqa_sample.jpg', 'rb')}
data = {'question': 'What is in the hand?'}

res = requests.post("http://localhost:8000/vqa", files=files, data=data)

print(res.json())
```

✳️ Next Modules You Can Link This Way

Task	Model	Endpoint
Object Detection	YOLOv8	/detect
Segmentation	U-Net / DINO-SAM	/segment
Image Classification	CLIP / ResNet	/classify
Bias Audit	Custom fairness report	/audit

Building: CV Lab Overview

A full-stack application that lets users interact with cutting-edge CV models like:

-  VQA (Visual Question Answering)
-  Object Detection (YOLOv8)
-  Image Segmentation (U-Net, SAM)
-  Bias Auditing Tools

All accessed via:

-  Web UI (React + Tailwind)
-  Backend APIs (FastAPI)
-  Serverless Deployment (Render, Vercel, or Azure Functions)
-  Real-time visualization (Plotly, OpenCV.js, TensorFlow.js)

Frontend Scaffold (React + Tailwind)

Directory Structure

cv-lab-frontend/

```
|  
|   public/  
|   |     images/  
|   src/  
|   |     components/  
|   |     |     VQADemo.jsx  
|   |     |     DetectionViewer.jsx  
|   |     |     SegmentationOverlay.jsx  
|   |     pages/  
|   |     |     Dashboard.jsx  
|   |     App.jsx  
|   |     index.js  
|   tailwind.config.js  
|   package.json
```

Sample Component (VQADemo.jsx)

```
import { useState } from "react";
import axios from "axios";
export default function VQADemo() {
  const [question, setQuestion] = useState("");
  const [answer, setAnswer] = useState("");
  const askQuestion = async () => {
    const formData = new FormData();
    formData.append("image", document.querySelector("#vqa-image").files[0]);
    formData.append("question", question);
    const res = await axios.post("https://api.example.com/vqa", formData);
    setAnswer(res.data.answer);
  };
  return (
    <div className="p-6">
      <input id="vqa-image" type="file" />
      <input
        type="text"
        value={question}
        onChange={(e) => setQuestion(e.target.value)}
        placeholder="Ask something..."
        className="border p-2"
      />
      <button onClick={askQuestion} className="bg-blue-500 text-white p-2">Ask</button>
      <p className="mt-4">Answer: {answer}</p>
    </div>
  );
}
```

Backend APIs (FastAPI + PyTorch/HuggingFace)

Live Endpoint Example

```
@app.post("/vqa")  
  
async def vqa(image: UploadFile, question: str = Form(...)):  
  
    image_bytes = await image.read()  
  
    result = answer_question(image_bytes, question)  
  
    return {"answer": result}
```

Extend with /detect, /segment, /audit, etc.

Deployment Guide: Serverless API Hosting

Options

- **Render:** Easy FastAPI deploy via GitHub
- **Vercel (frontend) + Supabase functions (backend)**
- **Azure Functions:** For scalability + identity management
- **HuggingFace Spaces:** For model hosting + UI together

Visualization Integration

In-browser Visuals

- `Plotly.js` for charts (bias metrics, loss curves)
- `OpenCV.js` for segmentation mask blending
- `TensorFlow.js` for real-time demos (e.g. webcam object detection)

Future Extensions

Feature	Tech Stack
Webcam object detection	TensorFlow.js + React hooks
U-Net mask overlay	Canvas + OpenCV.js
Bias audit report visualizer	Plotly bar charts & heatmaps
VQA image gallery	Cloud upload + MongoDB GridFS

Final Touches: Dashboard Design

Create a landing dashboard with cards for each demo:

```
const models = ["VQA", "Segmentation", "Detection", "Fairness Audit"];  
  
models.map((name) => (  
  
    <div className="bg-white p-4 rounded shadow">  
        <h2>{ name }</h2><button>Launch Demo</button>  
    </div>));
```

A complete full-stack scaffold for your **serverless Computer Vision Lab**, including:

- Web frontend (React + Tailwind)
- Backend API (FastAPI) with pre-trained model endpoints
- Deployment guides (for Render, Vercel, or HuggingFace Spaces)
- Real-time inference integration
- File structure and starter code snippets

Project Directory Structure

cv-lab/

```
├── backend/
│   ├── app.py
│   ├── models/
│   │   ├── vqa_model.py
│   │   ├── detection_model.py
│   │   └── segmentation_model.py
│   ├── requirements.txt
│   └── static/
│       └── sample.jpg
|
└── frontend/
    ├── public/
    └── src/
        ├── components/
        │   ├── VQADemo.jsx
        │   ├── ObjectDetection.jsx
        │   └── SegmentationMask.jsx
        ├── pages/Dashboard.jsx
        ├── App.jsx
        └── index.js
    └── tailwind.config.js
```

```
|   └── package.json
```

🧠 Backend API Scaffold (FastAPI)

```
🔧 app.py
```

```
from fastapi import FastAPI, UploadFile, Form

from models.vqa_model import answer_question

from models.detection_model import detect_objects

from models.segmentation_model import segment_image

app = FastAPI()

@app.post("/vqa")

async def vqa_endpoint(image: UploadFile, question: str = Form(...)):

    img_bytes = await image.read()

    result = answer_question(img_bytes, question)

    return {"answer": result}

@app.post("/detect")

async def detection_endpoint(image: UploadFile):

    img_bytes = await image.read()

    result = detect_objects(img_bytes)

    return {"boxes": result}

@app.post("/segment")

async def segmentation_endpoint(image: UploadFile):

    img_bytes = await image.read()

    mask = segment_image(img_bytes)

    return {"segmentation_mask": mask.tolist()}
```

```
🌐 vqa_model.py
```

```
from transformers import BlipProcessor, BlipForQuestionAnswering

from PIL import Image
```

```
import io

processor = BlipProcessor.from_pretrained("Salesforce/blip-vqa-base")

model = BlipForQuestionAnswering.from_pretrained("Salesforce/blip-vqa-base")

def answer_question(img_bytes, question):

    image = Image.open(io.BytesIO(img_bytes)).convert("RGB")

    inputs = processor(image, question, return_tensors="pt")

    out = model.generate(**inputs)

    return processor.decode(out[0], skip_special_tokens=True)
```

 Add similar modules for YOLO detection and U-Net segmentation in `detection_model.py` and `segmentation_model.py`.

Frontend UI Scaffold (React + Tailwind)

Setup

```
npx create-react-app frontend
```

```
cd frontend
```

```
npm install axios tailwindcss
```

```
npx tailwindcss init
```

Configure `tailwind.config.js` and add Tailwind directives to `src/index.css`.

VQADemo.jsx

```
import { useState } from "react";

import axios from "axios";

export default function VQADemo() {

  const [question, setQuestion] = useState("");

  const [answer, setAnswer] = useState("");

  async function handleSubmit() {

    const imageInput = document.getElementById("vqa-img");

    const formData = new FormData();

    formData.append("image", imageInput.files[0]);
```

```
        formData.append("question", question);

        const res = await axios.post("https://your-backend-url/vqa", formData);

        setAnswer(res.data.answer);

    }

    return (
        <div className="p-6 bg-white rounded shadow">
            <input id="vqa-img" type="file" />
            <input
                type="text"
                value={question}
                onChange={(e) => setQuestion(e.target.value)}
                placeholder="Ask your question..."/>
            className="border p-2 m-2"
        />
        <button onClick={handleSubmit} className="bg-blue-600 text-white px-4 py-2">
            Submit
        </button>
        <p className="mt-4">Answer: {answer}</p>
    </div>
);
}
```

-  Replicate this pattern for other modules (`ObjectDetection.jsx`, `SegmentationMask.jsx`) using their respective endpoints.

Deployment Guide

FastAPI Server

- Push `backend/` to GitHub
- Deploy to Render: Add a `start` command `uvicorn app:app --host 0.0.0.0 --port 10000`
- Set runtime: Python 3.10+, allow file uploads
- Add environment variables if needed

React Frontend

- Push `frontend/` to GitHub
- Deploy with Vercel or Netlify
- Point API endpoints to the deployed backend

Optional: Bundle frontend with HuggingFace Spaces for a single public-facing demo.

Add Live Visualization Libraries

- Use OpenCV.js for bounding box overlay
- Use TensorFlow.js for webcam-based demos
- Use Plotly.js for bias scores and training metrics
- Add canvas layers for segmentation masks